

Salary Prediction using SVM Regression

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

1. Data Pre-processing :

```
In [2]: df = pd.read_csv('Salary_dataset.csv')
```

```
In [3]: df.drop('Unnamed: 0', axis=1, inplace=True)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 612.0 bytes
```

```
In [5]: X = df.drop('Salary', axis=1)
Y = df['Salary']
```

```
In [6]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test = train_test_split(X, Y, test_size=0.20, random_state=42)
```

2. Create and Train SVM Regression Model :

```
In [7]: from sklearn.svm import SVR
```

```
In [8]: svr = SVR(kernel='linear')
svr.fit(x_train, y_train)
```

```
Out[8]: SVR
SVR(kernel='linear')
```

(<https://scikit-learn.org/1.4/modules/generated/sklearn.svm.SVR.html>)

3. Predict Test Set Results :

```
In [9]: y_pred = svr.predict(x_test)
```

```
In [10]: pd.DataFrame({'Actual':y_test, 'Predicted':y_pred})
```

Out[10]:

	Actual	Predicted
27	112636.0	62472.76
15	67939.0	62205.33
23	113813.0	62393.10
17	83089.0	62228.09
8	64446.0	62108.60
9	57190.0	62137.05

4. Evaluate Model Performance :

```
In [11]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [12]: mae = mean_absolute_error(y_test, y_pred)
mae
```

Out[12]: 22577.028333225255

```
In [13]: mse = mean_squared_error(y_test, y_pred)
mse
```

Out[13]: 943057673.9043975

```
In [14]: rmse = np.sqrt(mse)
rmse
```

Out[14]: 30709.244111576525

```
In [15]: r2 = r2_score(y_test, y_pred)
r2
```

Out[15]: -0.846251298098931


5. Visualization of Evaluate Model Performance

```
In [16]: # Plotting actual vs predicted
plt.figure(figsize=(10, 6))

# Plotting the actual values
plt.scatter(y_test, y_pred, color='blue', label='Actual vs Predicted')

# Plotting the line of best fit
plt.plot(y_test, y_test, color='red', label='Ideal line')

plt.title('Actual vs Predicted Salary')
plt.xlabel('Actual Salary')
plt.ylabel('Predicted Salary')
plt.legend()
plt.show()
```



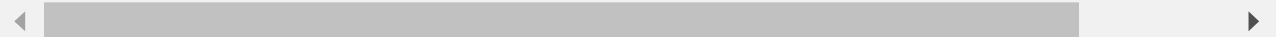
Universal bank Prediction

To build and train a model using Universal Bank records, and classify the customer is eligible to take Credit Card or not

```
In [17]: df = pd.read_csv('UniversalBank.csv')
df.head()
```

Out[17]:

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	Personal Loan	Securities Account	CD Account	Or
0	1	25	1	49	91107	4	1.6	1	0	0	1	0	
1	2	45	19	34	90089	3	1.5	1	0	0	1	0	
2	3	39	15	11	94720	1	1.0	1	0	0	0	0	
3	4	35	9	100	94112	1	2.7	2	0	0	0	0	
4	5	35	8	45	91330	4	1.0	2	0	0	0	0	



```
In [18]: df.isnull().sum()
```

```
Out[18]: ID                0
Age                0
Experience         0
Income            0
ZIP Code          0
Family            0
CCAvg             0
Education         0
Mortgage          0
Personal Loan     0
Securities Account 0
CD Account        0
Online            0
CreditCard       0
dtype: int64
```

```
In [19]: df.shape
```

Out[19]: (5000, 14)

```
In [20]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    5000 non-null   int64
1   Age                  5000 non-null   int64
2   Experience            5000 non-null   int64
3   Income               5000 non-null   int64
4   ZIP Code             5000 non-null   int64
5   Family               5000 non-null   int64
6   CCAvg                5000 non-null   float64
7   Education            5000 non-null   int64
8   Mortgage             5000 non-null   int64
9   Personal Loan        5000 non-null   int64
10  Securities Account    5000 non-null   int64
11  CD Account           5000 non-null   int64
12  Online               5000 non-null   int64
13  CreditCard           5000 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 547.0 KB
```

```
In [21]: df.describe()
```

Out[21]:

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	45.338400	20.104600	73.774200	93152.503000	2.396400	1.937913	1.881000
std	1443.520003	11.463166	11.467954	46.033729	2121.852197	1.147663	1.747666	0.839869
min	1.000000	23.000000	-3.000000	8.000000	9307.000000	1.000000	0.000000	1.000000
25%	1250.750000	35.000000	10.000000	39.000000	91911.000000	1.000000	0.700000	1.000000
50%	2500.500000	45.000000	20.000000	64.000000	93437.000000	2.000000	1.500000	2.000000
75%	3750.250000	55.000000	30.000000	98.000000	94608.000000	3.000000	2.500000	3.000000
max	5000.000000	67.000000	43.000000	224.000000	96651.000000	4.000000	10.000000	3.000000

```
In [22]: df.columns
```

Out[22]: Index(['ID', 'Age', 'Experience', 'Income', 'ZIP Code', 'Family', 'CCAvg', 'Education', 'Mortgage', 'Personal Loan', 'Securities Account', 'CD Account', 'Online', 'CreditCard'], dtype='object')

```
In [23]: df1 = df.drop(["ID", "ZIP Code"], axis = 1)
```


```
In [24]: df.head()
```

```
Out[24]:
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	Personal Loan	Securities Account	CD Account	Or
0	1	25	1	49	91107	4	1.6	1	0	0	1	0	
1	2	45	19	34	90089	3	1.5	1	0	0	1	0	
2	3	39	15	11	94720	1	1.0	1	0	0	0	0	
3	4	35	9	100	94112	1	2.7	2	0	0	0	0	
4	5	35	8	45	91330	4	1.0	2	0	0	0	0	

```
In [25]: plt.figure(figsize=(15,8))
plt.title("Showing Correlation between all the features", fontsize=20)
sns.heatmap(df1.corr(),annot = True, cmap='mako')
```

```
Out[25]: <Axes: title={'center': 'Showing Correlation between all the features'}>
```



```
In [26]: class_0 = df1[df1.CreditCard==0]
class_0.shape
```

```
Out[26]: (3530, 12)
```


```
In [27]: class_1 = df1[df1.CreditCard==1]
class_1.shape
```

```
Out[27]: (1470, 12)
```

```
In [28]: # Income vs Experience scatter plot
plt.xlabel('Income')
plt.ylabel('Experience')
plt.scatter(class_0['Income'],class_0['Experience'], color = 'green',marker='+')
plt.scatter(class_1['Income'], class_1['Experience'], color = 'red', marker='.')

```

```
Out[28]: <matplotlib.collections.PathCollection at 0x15f23ea2190>
```



```
In [29]: # Scaling the data using Standard Scaler
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled = scaler.fit(df1.drop('CreditCard',axis=1)).transform(df1.drop('CreditCard',axis=1))
df_scaled = pd.DataFrame(scaled, columns=df1.columns[:-1])
df_scaled.head()
```

```
Out[29]:
```

	Age	Experience	Income	Family	CCAvg	Education	Mortgage	Personal Loan	Securities Account	CD Account
0	-1.774417	-1.666078	-0.538229	1.397414	-0.193371	-1.049078	-0.555524	-0.325875	2.928915	-0.25354
1	-0.029524	-0.096330	-0.864109	0.525991	-0.250595	-1.049078	-0.555524	-0.325875	2.928915	-0.25354
2	-0.552992	-0.445163	-1.363793	-1.216855	-0.536720	-1.049078	-0.555524	-0.325875	-0.341423	-0.25354
3	-0.901970	-0.968413	0.569765	-1.216855	0.436103	0.141703	-0.555524	-0.325875	-0.341423	-0.25354
4	-0.901970	-1.055621	-0.625130	1.397414	-0.536720	0.141703	-0.555524	-0.325875	-0.341423	-0.25354

```
In [30]: # Splitting the columns in to dependent variable (x) and independent variable(y).
x = df_scaled
y = df1['CreditCard']
```

```
In [31]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

```
In [32]: # Apply SVM Model
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

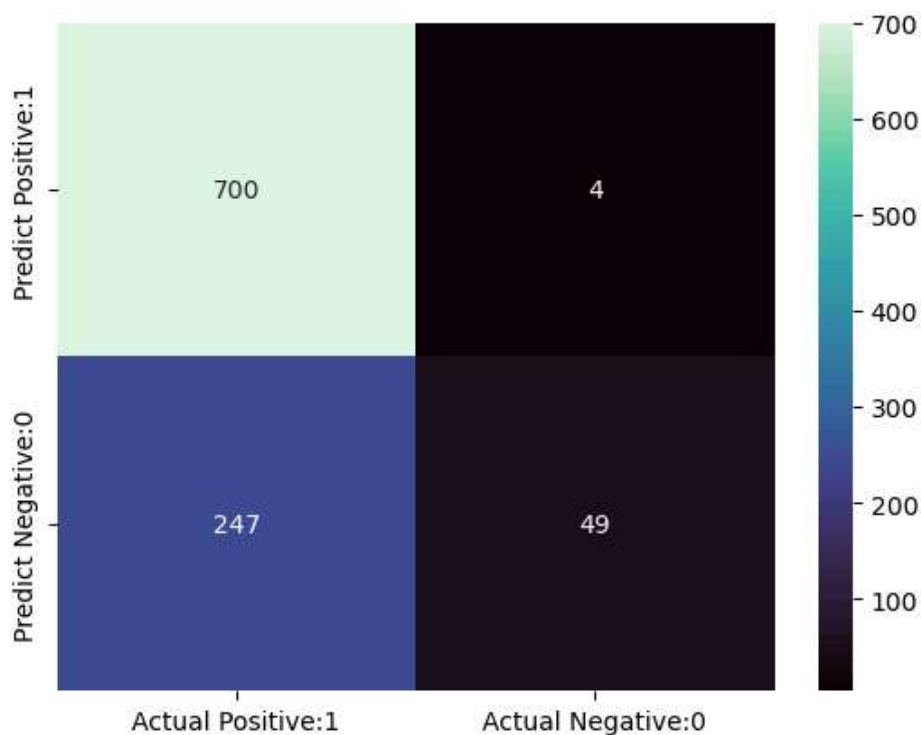
```
In [33]: svc=SVC()
svc.fit(x_train, y_train)
y_pred=svc.predict(x_test)
```

```
In [34]: print('Model accuracy : {0:0.3f}'. format(accuracy_score(y_test, y_pred)))
```

Model accuracy : 0.749

```
In [35]: # Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
index=['Predict Positive:1', 'Predict Negative:0'])
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='mako')
```

Out[35]: <Axes: >




```
In [36]: # Classification Report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.74	0.99	0.85	704
1	0.92	0.17	0.28	296
accuracy			0.75	1000
macro avg	0.83	0.58	0.56	1000
weighted avg	0.79	0.75	0.68	1000

SVM Kernel Functions

Sigmoid Kernel

```
In [37]: # Apply SVM model using Sigmoid Kernel function
Poly_svc=SVC(kernel='sigmoid', C=1).fit(x_train,y_train)
y_pred = Poly_svc.predict(x_test)
print('Model accuracy with rbf kernel : {0:0.3f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy with rbf kernel : 0.629

Linear Kernel

```
In [38]: # Apply SVM model using Linear Kernel function
linear_classifier=SVC(kernel='linear').fit(x_train,y_train)
y_pred = linear_classifier.predict(x_test)
print('Model accuracy with linear kernel : {0:0.3f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy with linear kernel : 0.747

Polynomial Kernel

```
In [39]: # Apply SVM model using Polynomial Kernel function
Poly_svc=SVC(kernel='poly', C=1).fit(x_train,y_train)
y_pred = Poly_svc.predict(x_test)
print('Model accuracy with rbf kernel : {0:0.3f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy with rbf kernel : 0.749

Gaussian RBF kernel

```
In [40]: # Apply SVM model using Gaussian RBF kernel function
rbf_svc=SVC(kernel='rbf').fit(x_train,y_train)
y_pred = rbf_svc.predict(x_test)
print('Model accuracy with rbf kernel : {0:0.3f}'.format(accuracy_score(y_test, y_pred)))
```

Model accuracy with rbf kernel : 0.749

House Price Prediction

```
In [41]: df=pd.read_csv("boston.csv")
```

```
In [42]: df.head()
```

Out[42]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33	36.2

```
In [43]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    int64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    int64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
In [44]: X = df.drop('MEDV', axis=1)
Y = df['MEDV']
```

```
In [45]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test = train_test_split(X, Y, test_size=0.20, random_state=42)
```

```
In [46]: from sklearn.svm import SVR
```

```
In [47]: svr = SVR(kernel='linear')
svr.fit(x_train, y_train)
```

```
Out[47]: SVR
SVR(kernel='linear')
(https://scikit-learn.org/1.4/modules/generated/sklearn.svm.SVR.html)
```

```
In [48]: y_pred = svr.predict(x_test)
```

```
In [49]: pd.DataFrame({'Actual':y_test, 'Predicted':y_pred})
```

```
Out[49]:
```

	Actual	Predicted
173	23.6	26.112730
274	32.4	32.551996
491	13.6	15.522970
72	22.8	24.140992
452	16.1	18.499756
...
412	17.9	-3.169195
436	9.6	13.189656
411	17.2	13.865112
86	22.5	21.442180
75	21.4	23.146530

102 rows × 2 columns

```
In [50]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [51]: mae = mean_absolute_error(y_test, y_pred)
mae
```

```
Out[51]: 3.1404227783347185
```

```
In [52]: mse = mean_squared_error(y_test, y_pred)
mse
```

```
Out[52]: 29.435701924289845
```

```
In [53]: rmse = np.sqrt(mse)
rmse
```

```
Out[53]: 5.4254678991115455
```

```
In [54]: r2 = r2_score(y_test, y_pred)
r2
```

```
Out[54]: 0.5986065268181071
```

```
In [ ]:
```