

# Letter Recognition

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df = pd.read_csv("letter-recognition.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	letter	xbox	ybox	width	height	onpix	xbar	ybar	x2bar	y2bar	xybar	x2ybar	xy2bar	xedge	xedgey	yedge	y
0	T	2	8	3	5	1	8	13	0	6	6	10	8	0	8	0	
1	I	5	12	3	7	2	10	5	5	4	13	3	9	2	8	4	
2	D	4	11	6	8	6	10	6	2	6	10	3	7	3	7	3	
3	N	7	11	6	6	3	5	9	4	6	4	4	10	6	10	2	
4	G	2	1	3	1	1	8	6	6	6	6	5	9	1	7	5	

```
In [4]: df.describe()
```

```
Out[4]:
```

	xbox	ybox	width	height	onpix	xbar	ybar	x2bar
count	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000
mean	4.023550	7.035500	5.121850	5.37245	3.505850	6.897600	7.500450	4.628600
std	1.913212	3.304555	2.014573	2.26139	2.190458	2.026035	2.325354	2.699968
min	0.000000	0.000000	0.000000	0.00000	0.000000	0.000000	0.000000	0.000000
25%	3.000000	5.000000	4.000000	4.00000	2.000000	6.000000	6.000000	3.000000
50%	4.000000	7.000000	5.000000	6.00000	3.000000	7.000000	7.000000	4.000000
75%	5.000000	9.000000	6.000000	7.00000	5.000000	8.000000	9.000000	6.000000
max	15.000000	15.000000	15.000000	15.00000	15.000000	15.000000	15.000000	15.000000

```
In [5]: df.shape
```

```
Out[5]: (20000, 17)
```

```
In [6]: df.columns
```

```
Out[6]: Index(['letter', 'xbox ', 'ybox ', 'width ', 'height', 'onpix ', 'xbar ',
'ybar ', 'x2bar', 'y2bar ', 'xybar ', 'x2ybar', 'xy2bar', 'xedge ',
'xedgey', 'yedge ', 'yedgey'],
dtype='object')
```

```
In [7]: # Let's 'reindex' the column names
df.columns = ['letter', 'xbox', 'ybox', 'width', 'height', 'onpix', 'xbar',
              'ybar', 'x2bar', 'y2bar', 'xybar', 'x2ybar', 'xy2bar', 'xedge',
              'xedgey', 'yedge', 'yedgey']
df.columns
```


```
Out[7]: Index(['letter', 'xbox', 'ybox', 'width', 'height', 'onpix', 'xbar', 'ybar',
              'x2bar', 'y2bar', 'xybar', 'x2ybar', 'xy2bar', 'xedge', 'xedgey',
              'yedge', 'yedgey'],
              dtype='object')
```

```
In [8]: order = list(np.sort(df['letter'].unique()))
print(order)
```

```
['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S',
 'T', 'U', 'V', 'W', 'X', 'Y', 'Z']
```

```
In [9]: plt.figure(figsize=(16, 8))
sns.barplot(x='letter', y='xbox', data=df, order=order)
```

```
Out[9]: <Axes: xlabel='letter', ylabel='xbox'>
```




```
In [10]: df1 = df.groupby('letter').mean()
df1.head()
```

```
Out[10]:
```


	xbox	ybox	width	height	onpix	xbar	ybar	x2bar	y2bar	xybar	x2ybar	
letter												
A	3.337136	6.975919	5.128010	5.178707	2.991128	8.851711	3.631179	2.755387	2.043093	7.802281	2.338403	8
B	3.985640	6.962141	5.088773	5.169713	4.596606	7.671018	7.062663	5.366841	5.571802	7.954308	5.506527	6
C	4.031250	7.063859	4.701087	5.296196	2.775815	5.437500	7.627717	5.927989	7.177989	8.773098	7.494565	11
D	4.023602	7.244720	5.170186	5.288199	4.026087	7.539130	6.806211	5.921739	6.508075	8.166460	5.111801	5
E	3.727865	6.944010	4.756510	5.201823	3.679688	5.966146	7.352865	4.223958	7.585938	8.507812	6.242188	10

```
In [11]: plt.figure(figsize=(18, 10))
sns.heatmap(df1)
```

```
Out[11]: <Axes: ylabel='letter'>
```



```
In [12]: plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='letter', order=df['letter'].value_counts().index)
plt.title("Distribution of Letters")
plt.show()
```



### Data Preparation

```
In [13]: # average feature values
round(df.drop('letter', axis=1).mean(), 2)
```

```
Out[13]: xbox      4.02
ybox      7.04
width     5.12
height    5.37
onpix     3.51
xbar      6.90
ybar      7.50
x2bar     4.63
y2bar     5.18
xybar     8.28
x2ybar    6.45
xy2bar    7.93
xedge     3.05
xedgey    8.34
yedge     3.69
yedgey    7.80
dtype: float64
```

```
In [14]: # splitting into X and y
X = df.drop("letter", axis = 1)
y = df['letter']
```

```
In [15]: # Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Splitting into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

## Linear SVM

```
In [16]: # Create and train the Linear SVM model
svm_model = SVC(kernel='linear', C=1)
svm_model.fit(X_train, y_train)

# predict
y_pred = svm_model.predict(X_test)
```

```
In [17]: # Calculate accuracy and display results
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
```


Accuracy: 0.86

```
In [18]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
A	0.87	0.94	0.90	149
B	0.81	0.89	0.85	153
C	0.89	0.85	0.87	137
D	0.76	0.92	0.83	156
E	0.83	0.91	0.87	141
F	0.80	0.90	0.85	140
G	0.75	0.81	0.78	160
H	0.65	0.58	0.61	144
I	0.90	0.87	0.89	146
J	0.86	0.87	0.87	149
K	0.75	0.78	0.77	130
L	0.95	0.86	0.91	155
M	0.93	0.95	0.94	168
N	0.96	0.90	0.93	151
O	0.90	0.77	0.83	145
P	0.96	0.84	0.90	173
Q	0.84	0.80	0.82	166
R	0.75	0.84	0.79	160
S	0.76	0.73	0.74	171
T	0.91	0.90	0.91	163
U	0.94	0.90	0.92	183
V	0.91	0.90	0.90	158
W	0.90	0.96	0.93	148
X	0.93	0.90	0.92	154
Y	0.95	0.88	0.91	168
Z	0.88	0.82	0.85	132
accuracy			0.86	4000
macro avg	0.86	0.86	0.86	4000
weighted avg	0.86	0.86	0.86	4000

```
In [19]: # Create a confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix
plt.figure(figsize=(12, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



## Non-Linear SVM

```
In [20]: # Train a non-linear SVM model (e.g., using a radial basis function kernel)
svm = SVC(kernel='rbf', C=1.0, gamma='scale', random_state=101)
svm.fit(X_train, y_train)
```

```
Out[20]: SVC
(https://scikit-learn.org/1.4/modules/generated/sklearn.svm.SVC.html)
SVC(random_state=101)
```

```
In [21]: # Make predictions on the test set
y_pred = svm.predict(X_test)
```

```
In [22]: # Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.95


Classification Report:

	precision	recall	f1-score	support
A	0.98	0.99	0.99	149
B	0.87	0.96	0.91	153
C	0.98	0.92	0.95	137
D	0.88	0.98	0.93	156
E	0.91	0.96	0.93	141
F	0.90	0.96	0.93	140
G	0.92	0.96	0.94	160
H	0.97	0.82	0.89	144
I	0.98	0.93	0.95	146
J	0.94	0.96	0.95	149
K	0.91	0.90	0.91	130
L	0.99	0.93	0.96	155
M	0.98	0.99	0.99	168
N	0.98	0.93	0.95	151
O	0.92	0.94	0.93	145
P	0.99	0.89	0.94	173
Q	0.99	0.98	0.98	166
R	0.83	0.94	0.88	160
S	0.99	0.98	0.98	171
T	0.98	0.95	0.97	163
U	0.97	0.96	0.96	183
V	0.99	0.95	0.97	158
W	0.96	0.99	0.97	148
X	0.96	0.99	0.97	154
Y	0.99	0.99	0.99	168
Z	0.98	0.96	0.97	132
accuracy			0.95	4000
macro avg	0.95	0.95	0.95	4000
weighted avg	0.95	0.95	0.95	4000

In [23]:

```
# Create a confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Visualize the confusion matrix using a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='pink')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



## Social Network Ads Prediction using SVM

In [24]: `df = pd.read_csv('Social_Network_Ads.csv')`  
`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```



```
In [25]: df.head()
```

```
Out[25]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [26]: # Assuming your DataFrame is named 'df'  
df.drop(['User ID', 'Gender'], axis=1, inplace=True)
```

```
In [27]: df.head()
```

```
Out[27]:
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

```
In [28]: X = df.drop('Purchased', axis=1)  
Y = df['Purchased']
```

```
In [29]: from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test = train_test_split(X, Y, test_size=0.20, random_state=42)
```

```
In [30]: from sklearn.svm import SVC
```

```
In [31]: svm_clf = SVC(kernel='rbf', C=1.0)  
svm_clf.fit(x_train, y_train)
```

```
Out[31]:
```

▼ SVC ⓘ ?

(<https://scikit-learn.org/1.4/modules/generated/sklearn.svm.SVC.html>)  
SVC()

```
In [32]: y_pred = svm_clf.predict(x_test)
```

```
In [33]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [34]: score = accuracy_score(y_test, y_pred)  
print(score)
```

0.7375

```
In [35]: mat = confusion_matrix(y_test, y_pred)
print(mat)
```

```
[[49  3]
 [18 10]]
```


```
In [36]: report = classification_report(y_test, y_pred)
print(report)
```

	precision	recall	f1-score	support
0	0.73	0.94	0.82	52
1	0.77	0.36	0.49	28
accuracy			0.74	80
macro avg	0.75	0.65	0.66	80
weighted avg	0.74	0.74	0.71	80

```
In [37]: from sklearn.inspection import DecisionBoundaryDisplay
```

```
In [38]: # Plot Decision Boundary
DecisionBoundaryDisplay.from_estimator(
    svm_clf,
    x_train,
    response_method="predict",
    cmap=plt.cm.Spectral,
    alpha=0.8,
    xlabel='Age',
    ylabel='EstimatedSalary',
)

# Scatter plot
plt.scatter(x_train['Age'], x_train['EstimatedSalary'],
            c=y_train,
            s=20, edgecolors="k")
plt.show()
```



## IRIS Flower Prediction using Non-Linear SVM

```
In [39]: df=pd.read_csv("IRIS.csv")
```

```
In [40]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [41]: df.species.unique()
```

```
Out[41]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

```
In [42]: df['species'] = df['species'].replace({'Iris-setosa':1, 'Iris-versicolor':2, 'Iris-virginica':3})
```

```
In [43]: df.head()
```

```
Out[43]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	1
1	4.9	3.0	1.4	0.2	1
2	4.7	3.2	1.3	0.2	1
3	4.6	3.1	1.5	0.2	1
4	5.0	3.6	1.4	0.2	1

```
In [44]: # Select X and y without feature_names
```

```
X = df[['sepal_length','petal_width']].values # Select the first two features for visualization  
y = df.species
```

```
In [45]: from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

```
In [46]: from sklearn.svm import SVC
```

```
In [47]: svm_clf = SVC(kernel='poly',degree=5, random_state=42)  
svm_clf.fit(x_train, y_train)
```

```
Out[47]:
```

SVC

<https://scikit-learn.org/1.4/modules/generated/sklearn.svm.SVC.html>

SVC(degree=5, kernel='poly', random\_state=42)

```
In [48]: y_pred = svm_clf.predict(x_test)
```

```
In [49]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [50]: score = accuracy_score(y_test, y_pred)  
print(score)
```

```
1.0
```

```
In [51]: report = classification_report(y_test, y_pred)
print(report)
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	10
2	1.00	1.00	1.00	9
3	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [52]: mat = confusion_matrix(y_test, y_pred)
print(mat)
```

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

```
In [53]: # Define Grid for Visualization
```


```
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02),
                     np.arange(y_min, y_max, 0.02))
```

```
In [54]: # Make Predictions on Mesh Grid
```

```
Z = svm_clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
```

```
In [55]: # Plot Decision Boundary and Data Points
```

```
plt.figure(figsize=(8, 6))
plt.contourf(xx, yy, Z, alpha=0.8)
plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', marker='o')
plt.xlabel('Sepal Length')
plt.ylabel('Sepal Width')
plt.title('Non-Linear SVM Decision Boundary on Iris Dataset')
plt.show()
```



```
In [ ]:
```